of existence; fortunately, less-threatening parts of the physical world are similarly real—a change in my thinking, by *itself*, cannot make my body, my friends, or my neighborhood go out of existence, nor thankfully can a change in anyone else's. To bring about a change in the world, you have to do more than just think about it. However, if we want to maintain a notion of independent reality, we should consider to what extent the research on social construction challenges the idea. If a strong case can be made for the claim that *reality* is socially constructed and, further, that what's socially constructed is *not* independently real, then we may have to consider a more radically revisionary view about the world.

My project in this paper is to explore the claim that *reality* is socially constructed; more broadly, I hope to show how debates over such philosophical notions as "truth," "knowledge," and "reality" can be relevant to feminist and antiracist politics. In the following section I will consider what it means to say that something is socially constructed and will distinguish several senses of the term (allowing that there are also many others); I've chosen to set out this rather complex set of distinctions because their differences become significant in the arguments that follow. I'll then turn to consider how far the claim that reality is socially constructed commits us to denying that the world is, at least in part, independent of us. I will examine a strategy of argument claiming that because *knowledge* is socially constructed, there is no objective (and so no independent) reality. I argue, however, that even if this strategy provides good reason for rejecting one conception of "objective reality," this does not force us into either skepticism or idealism, for there are other ways of conceiving what it means to be real and other ways of conceiving an "independent" reality. My intention here is not to offer an argument *for* realism, or *for* an independent reality; rather, it is (more modestly) to understand and evaluate some of the arguments that may seem to challenge such commitments.

2. Social Construction

As mentioned above, the notion of "social construction" is applied to a wide variety of items and seemingly with rather different senses.[7] At least initially it is useful to think of social constructions on the model of artifacts.[8] In addition to straightforward artifacts like washing machines and power drills, there is a clear sense in

[7] A good place to begin in considering recent uses of the phrase 'social construction' is Berger and Luckmann, op. cit.; for a recent survey of uses of the term, especially in the sociology of knowledge, see Sergio Sismondo, "Some Social Constructions," *Social Studies Science* 23 (1993): 515–53.

[8] Whether, ultimately, social construction should be understood in terms of artifacts is controversial; in particular, the suggestions that artifacts require agents to produce them and that artifacts invite a matter-form analysis have been targeted as problematic. See, e.g., Judith Butler, *Bodies That Matter* (New York: Routledge, 1993), ch. 1.

which, for example, the Supreme Court of the United States and chess games are artifacts, as are languages, literature, and scientific inquiry. Because each of these depends for its existence on a complex social context, each is in the broad sense in question a social construction. So, let's say:

> *Generic social construction*: Something is a social construction in the generic sense just in case it is an intended or unintended product of a social practice.

Although it is fair to say that, generally speaking, social constructions are artifacts, this leaves much open, since there are many different kinds of artifacts and ways of being an artifact. In perhaps the paradigm case of artifacts, human beings play a causal role in bringing an object into existence in accordance with a design plan or to fulfill a specific function. However, the idea of artifact, and with it the idea of social construction, extends well beyond this paradigm case: Human intention or design is not always required (natural languages and cities are certainly artifacts, but they are not the work of an intentional agent or artisan); in other cases the issue does not concern origins but whether the conditions for being the kind of object in question make reference to social practices. For example, categories of individuals such as professors or wives and other *social* kinds count as social constructions because the conditions for being a member of the kind or category include *social* (properties and) relations: The category of wives counts because you can't be a wife unless you are part of a social network that provides for an institution of marriage.

These examples suggest a distinction between *causal* and *constitutive* senses of construction that is important, for it makes a big difference to how we should evaluate the claim that something is socially constructed. For example, in some contexts, to say that "gender" is socially constructed is to make a claim about the causes of gender-coded traits in individuals; that is, it is to claim that insofar as women are feminine and men are masculine, this is due (at least in part) to social causes and is not biologically determined.[9] Presumably, in order to evaluate this claim, we would have reason to consider data from the social sciences, including psychology, sociology, anthropology, and history.

However, in other contexts the claim that "gender" is socially constructed is not a causal claim; rather, the point is constitutive: Gender should be understood as a social category whose definition makes reference to a broad network of

[9] In the context of psychology, the claim is often more specific, viz., that our own self-attributed gender, or gendered sense of self, is a result of social forces and is not biologically determined. In the feminist literature, one finds the term 'gender identity' being used in different ways, e.g., sometimes for the psychological phenomenon of self-ascribed gender, sometimes for one's gendered characteristics more broadly, and sometimes for the social category one has been ascribed to.

social relations, and it is not simply a matter of anatomical differences.[10] In this case, gender is introduced as an analytical tool to explain a range of social phenomena, and we evaluate the claim by considering the theoretical usefulness of such a category.[11] There is room for much debate here, not only over the question whether we should employ such a category, but if we do, how we should define it, that is, what social relations (or clusters of social relations) constitute the groups *men* and *women*. (The debates here parallel others in social theory: One might debate whether the category "middle class" is useful to explain a range of social phenomena and, if so, how we should define it.)

To help keep distinct these different ways in which the social can function in construction, let's distinguish:

> *Causal construction*: Something is causally constructed iff social factors play a causal role in bringing it into existence or, to some substantial extent, in its being the way it is.
>
> *Constitutive construction*: Something is constitutively constructed iff in defining it we must make reference to social factors.[12]

We need to consider causal construction further, for things get quite complicated when we consider how social factors can have an effect on the world (we'll also return to constitutive construction below). At least in the case of human beings, the mere fact of how we are (even potentially) described or classified can have a direct impact on our self-understandings and our actions, because typically these descriptions and classifications bring with them normative expectations

[10] A significant amount of feminist work over the past two decades has been concerned with gender as a social category. For important examples, see Joan Scott, "Gender: A Useful Category of Historical Analysis," *American Historical Review* 91 (1986): 1053–75; Donna Haraway, "'Gender' for a Marxist Dictionary," in her *Simians, Cyborgs, and Women*. See also, Sally Haslanger, "On Being Objective and Being Objectified," chapter 1 of this volume, esp. secs. 2 and 4; and Linda Alcoff, "Cultural Feminism versus Post-Structuralism," *Signs* 13 (1988): 405–36. esp. 433–36.

[11] A "social constructivist" approach to a given domain is typically set in opposition to an "essentialist" approach. Because there are different senses in which things are socially constructed, it is not surprising that the term 'essentialism' also has a variety of uses. For example, those who are concerned to assert social construction as a causal thesis about the social origins of certain traits or capacities tend to interpret the opposing essentialism as a commitment to biological determinism; however, in contexts where constructionists are postulating a social category, essentialism is usually taken to be the view that all members of the category share some (intrinsic?) feature(s). Both of these "essentialisms" are different from the kind of modal (or Aristotelian) essentialism discussed in contemporary analytic metaphysics.

[12] I intend this definition of constitutive construction to be applicable to objects, kinds of objects. properties, or concepts. Some may find it puzzling that I speak of defining objects and features of objects, since contemporary philosophers have often insisted that terms or concepts are the (only) proper subjects of definition. I take it, however, that the seeming unanimity on this point is breaking down, and I want to allow broad flexibility in the notion of constitutive construction. See Kit Fine, "Essence and Modality," *Philosophical Perspectives* 8 (1994): 1–16.

and evaluations. This works in several ways. Forms of description or classification provide for kinds of intention; for example, given the classification "cool," I can set out to become cool, or avoid being cool, and so on. But also, such classifications can function in justifying behavior—for example, "we didn't invite him, because he's not cool"—and such justifications, in turn, can reinforce the distinction between those who are cool and those who are uncool.[13]

The main point to note here is that our classificatory schemes, at least in social contexts, may do more than just map preexisting groups of individuals; rather our attributions have the power to both establish and reinforce groupings which may eventually come to "fit" the classifications. In such cases, classificatory schemes function more like a script than a map. This gives us a narrower conception of social construction falling under the more general rubric of causal construction. On this conception something is socially constructed if what or how it is depends on a kind of feedback loop involving activities such as naming or classifying.[14] Sometimes this form of construction is called "linguistic" or "discursive" construction,[15] so I'll keep with this terminology:

> *Discursive construction*: Something is discursively constructed just in case it is the way it is, to some substantial extent, because of what is attributed (and/or self-attributed) to it.

I'd say that there is no doubt that in this sense you and I are socially constructed: We are the individuals we are today at least partly as a result of what has been attributed (and self-attributed) to us. In other words, there is a sense in which adult human beings are a special kind of artifact.[16]

[13] No doubt with time the term 'cool' will come to seem awkward and dated as an honorific; such changes in the terminology used to establish social groups are inevitable. If the reader finds "being cool" no longer socially desirable, substitute in the examples whatever term currently functions in its place.

[14] Ian Hacking explicitly mentions such a "feedback loop" in "The Sociology of Knowledge about Child Abuse," *Nous* 22 (1988): 55. See also Ian Hacking, "The Making and Molding of Child Abuse," *Critical Inquiry* 17 (1990/91): 253–88.

[15] Note that a discourse, and so discursive construction, will involve more than spoken language. See Nancy Fraser, "The Uses and Abuses of French Discourse Theories for Feminist Politics," in Nancy Fraser and Sandra L. Bartky, eds., *Revaluing French Feminism* (Bloomington: Indiana University Press, 1992). 177–94. For a clear explanation of one feminist appropriation of Foucault's notion of "discourse," see, also, Joan Scott, "Deconstructing Equality-verses-Difference: or, The Uses of Poststructuralist Theory for Feminism," *Feminist Studies* 14 (1988): 33–50.

[16] The claim that adult human beings are artifacts allows that we are constructed from "natural" materials—e.g., flesh and blood—though this would seem to presuppose a clear distinction between the natural and the social. For helpful discussions questioning feminist uses of this distinction, see Marilyn Frye, *The Politics of Reality* (Freedom, Calif.: The Crossing Press, 1983). 34–37; Moira Gatens, "A Critique of the Sex-Gender Distinction," in J. Allen and P. Patton, eds., *Beyond Marxism? Interventions after Marx* (Sydney: Interventions Publications, 1983); and Judith Butler, *Bodies That Matter*, esp. ch. 1.

Things get even more complicated, though, because there's still another wrinkle to consider. The idea of discursive construction depends on there being descriptions, distinctions, and classifications at hand whose attribution to things makes a difference—I am the way I am today because people have had the linguistic and conceptual resources to describe me as, for example, "smart" or "stupid," "attractive" or "ugly." There is yet another sense of social construction in which it makes sense to say that *these classificatory schemes themselves*—our distinctions such as smart or stupid, attractive or ugly, rather than the things that respond to them—are socially constructed. Very roughly, to say that such a scheme is socially constructed is to say that its use is determined, not by the "intrinsic" or "objective" features of the objects to which it is applied, but by social factors.[17]

This characterization is purposely vague; so to help us explore some of the issues involved in it let's go back to the example of "being cool": In considering our use of the distinction between those who are cool and those who are uncool, it is plausible to conclude that the distinction is not capturing intrinsic differences between people; rather it is a distinction marking certain social relations—that is, it distinguishes status in the in-group—and the fact that it is employed in any given context is a reflection of the importance of in-group and out-group relations. For example, suppose I need a way to establish a cohort; I do so by calling those I like "cool" and those I don't "uncool." The distinction does not capture a difference in the individuals so-called except insofar as they are related to me (based on my likes and dislikes), and its use in the context is determined not by the intrinsic or objective coolness of the individuals but by the social task of establishing a cohort.[18]

[17] Note that I do not mean to equate "intrinsic" with "objective" features. Intrinsic features are, roughly, those that an object has "by virtue of itself alone": i.e., they are nonrelational properties. In this context we can take objective features to be (very roughly) those that an object has independently of its representation by an inquiring subject. At the very least. not all objective features are intrinsic; e.g., *orbiting the sun* is an extrinsic or relational property of the earth that is plausibly objective. The earth orbits the sun or doesn't, independently of what we think about it. Note that, in the examples below, I assume that in attributing "coolness" to someone we are suggesting that they have an *intrinsic* quality of coolness. I'm now not convinced that this is right, because the attribution of coolness seems to allow that there are some relational features relevant to someone's being cool, e.g., owning an electric guitar; instead it may be that what's at issue is the "objectivity" of coolness—we suggest that coolness has nothing to do with our representations, when in fact it does. For ease of presentation, I've focused on intrinsicness in the examples, though a more complete discussion would have to show greater sensitivity to the particular nuances of "cool."

[18] Though if I am successful and there is solidarity in the cohort, we may come to act alike, dress alike, value similar things, etc., and this can provide substantive content to the notion of "coolness": eventually there may be a genuine (intrinsic) difference between the cool and uncool. For an up-to-date sample of who's and what's *really* cool, you can contact the "Who's Cool in America Project" though the World Wide Web at: http://www.attisv.com/-getconl/index.shtml. You can even submit an application stating why you think you are cool, and a "CoolBoard" determines whether you are.

Noting the influence of social forces upon the distinctions we draw, let us define this third form of social construction, as follows:

> *Pragmatic construction*: A classificatory apparatus (be it a full-blown classification scheme or just a conceptual distinction or descriptive term) is socially constructed just in case its use is determined, at least in part, by social factors.

Construed in its weakest form, the point in claiming that a given distinction is pragmatically constructed is simply to say that our use of that distinction is as much due to contingent historical and cultural influences as to anything else; we inherit vocabularies and classificatory projects and decide between alternatives based on utility, simplicity, and such. This point is easy to grant; it would be hard to deny that the discursive resources we employ are socially conditioned in these ways and more. In a stronger form, however, the point is that social factors *alone* determine our use of the distinction in question; in short it is to emphasize that there's no "fact of the matter" that the distinction captures. So let's distinguish two kinds of pragmatic construction:

> A distinction is *weakly pragmatically constructed* if social factors only partly determine our use of it.
>
> A distinction is *strongly pragmatically constructed* if social factors wholly determine our use of it, and it fails to represent accurately any "fact of the matter."[19]

We'll come back to the weak form of pragmatic construction shortly; let me first unpack this strong form further, because there is an ambiguity in the suggestion that there's no fact of the matter that such a pragmatically constructed distinction captures. In the example of "cool," I use the term to establish my cohort, and in doing so my ascriptions are guided by my likes and dislikes; so there may be a real social distinction (admittedly parochial) that corresponds to my use—I call Mary

[19] As will become clear in what follows, this characterization of strong pragmatic construction is oversimplified, even misleading, in its suggestion that in the relevant cases social concerns *wholly* determine our use of the distinction. In the case of "cool," because our use tracks *some* real distinction (status with respect to the in-group), facts about individuals we label "cool dudes" matter to whether we apply the term (e.g., whether they dress a certain way, behave a certain way, etc.). What I'm trying to capture, however, is the fact that our usage is not being guided by some actual property of individuals that corresponds to the intended content (intrinsic coolness) and that what property (or properties) substitutes for it is determined by social concerns. In terms I introduce below, we could say that what operative concept substitutes for the manifest concept is determined wholly by social factors. Let me also emphasize that by labeling the different forms of pragmatic construction "weak" and "strong" I do not intend to imply that the more socially motivated a distinction, the less real; we have very strong social reasons for marking certain real distinctions.

and George "cool," Susan and John "uncool," and the application of the terms corresponds to who I like and who I don't. But note also that in attributing "coolness" to someone, I'm doing so with the background assumption in play that the "coolness" is an intrinsic feature of the individual and is not merely a matter of whom I like. In calling Mary and George "cool," I'm suggesting that there is something cool *about them* that has nothing to do with me—supposedly, it's *their coolness* that warrants my use of the term. It is here that the question of fact arises: Insofar as I am attributing intrinsic coolness to someone, my attribution misfires since no one is, so to speak, cool *in themselves*. In such cases I want to say that my attributions of coolness are false—there is no fact about their coolness that I am accurately representing, even if my use of the terms corresponds to some other features of the individuals, for example, whether or not I like them.[20] So, *strong pragmatic constructions are, in an important sense, illusions projected onto the world; their use might nevertheless track—without accurately representing—a genuine distinction*. The main point is that in cases of strong pragmatic construction there are no available facts corresponding to the intended content—in the case at hand, about intrinsic coolness or uncoolness—that my attributions could be tracking, so instead, we might conclude, they must be functioning *wholly* as a means to a social goal.

On the face of it, there is a significant difference between weak and strong pragmatic construction. In cases of weak pragmatic construction our choices of descriptive terms, classificatory schemes, and so on, are conditioned by social factors (values, interests, history, etc.), but of course this is compatible with those terms' and classifications' capturing real facts and distinctions. The world provides us with more facts and distinctions than we could ever know what to do with; acknowledging that what ones we bother to notice or name is largely determined by our background and interests does not impugn in any general way the accuracy of our attributions.[21] In cases of strong pragmatic construction, however, the attributions are, by hypothesis, not accurately capturing facts, though there is an illusion that they are.

It is important to note that because in the case of pragmatic construction, what's constructed is (at least primarily) a distinction or classificatory scheme, the thought that our classifications are socially constructed leads naturally to the idea that *knowledge* is socially constructed. Given the preceding discussion, we must allow that there are different ways to cash out the claim that knowledge

[20] Others will likely question whether I've made a genuine assertion at all, and still others may suggest that I've said something true but misleading. The issues that arise here parallel debates over antirealism and realism in other domains; in effect I am proposing here an "error theory" about coolness, where others might endorse a realism or noncognitivism (though shortly I will modify this view somewhat). But this is a debate we need not settle here. For a general discussion of the alternatives, see Geoffrey Sayre-McCord, "Introduction: The Many Moral Realisms," in his *Essays on Moral Realism* (Ithaca, N.Y.: Cornell University Press, 1988); and Paul Boghossian, "The Status of Content," *Philosophical Review* 99 (1990): 157–84.

[21] See John Dupré, *The Disorder of Things* (Cambridge, Mass.: Harvard University Press, 1993), ch. 1.

is socially constructed, but we can cast two of them in terms of weak and strong pragmatic construction. Roughly:

> Our *knowledge* is weakly/strongly socially constructed (in the relevant senses) iff the distinctions and classifications we employ in making knowledge claims are weakly/strongly pragmatically constructed.

We now have three basic senses of construction to work with: causal, constitutive, and pragmatic. To see how these can become intertwined let's consider the project of debunking strong pragmatic constructions. Return once again to the case of "cool": Attributions of "coolness" have an effect on how individuals interact. "Cool dudes" are *discursively constructed*. But on the analysis I've been proposing, this happens as a result of a false and importantly misleading representation of the facts. I am suggesting that in contexts where "coolness" functions as a serious form of evaluation, there is general complicity in the belief that cool behavior is a result of a character trait (the person's *being* cool) that is the real basis for the evaluation. Cool dudes want their coolness, so to speak, to "shine through" in their behavior, dress, and so on, so that they will win approval by the in-group; and the in-group acknowledges a distinction between *being* cool and just acting cool. Cool things (objects, dress, actions) are the things cool people approve of (or would approve of). To debunk the belief that there is a special quality of coolness that warrants the designation "cool," we show that there is no such property of "coolness" (so understood) and, in fact, that the application of the term "cool" is determined wholly by the interests and concerns of the in-group. In other words, "coolness," when debunked, is revealed as a *constitutive construction*; that is, the concept doing the work of determining when the term should be applied makes essential reference to social factors (that is, in-group status).

But we must be careful here: What counts as the concept "cool"? Once we have disrupted the coolness illusion, there seem to be two different concepts playing a role in our use of the term. On the one hand, there is the concept that actually determines how we apply the term to cases, that is, (roughly) being such as to conform to the standards of the in-group. Let's call this the *operative* concept. On the other hand, there is the concept that users of the term typically take (or took) themselves to be applying, that is, being intrinsically or objectively cool, where this is supposed to be the objective basis for the in-group standards. Let's call this the *manifest* concept. In attributing "coolness" (or "uncoolness") to someone, we are using the apparent objectivity of the manifest concept of "coolness" as a mask for the explicitly social content of the operative concept. But which of these two concepts is the concept "cool"? Both seem to be reasonable candidates: When we sincerely say that someone is "cool," or when we begin the debunking project by insisting that we are mistaken in our attributions of coolness—no one is *really* cool—what's at issue is the manifest concept; but once the debunking

project has taken hold, it is tempting to break the illusion by saying that we were wrong about what "coolness" involved and that coolness itself is a constitutive construction. In this we shift from thinking of "cool" in terms of the manifest concept to the operative concept.

So in saying that "coolness" is a social construction, one could have in mind either (or both) (i) that "cool" individuals are discursively constructed (the pattern of behavior found in "cool" individuals is caused by a complex system of attribution and response) or (ii) that the operative concept expressed when we use the term "cool" is constitutively constructed (our use of the term "cool" is actually governed by conditions that concern in-group status, and the content normally associated with the term is a mask for these social conditions). These two ideas are intertwined because the discursive construction of "cool" individuals partly depends upon the (masked) attribution of the constitutively constructed concept "cool."

To see more clearly how these different kinds of social construction function, let's shift from the somewhat artificial example of "cool" I've been using to something more substantive and, for some, more familiar. I'll run briefly through the different kinds by using an example of the social construction of gender. As usual, allow at least a provisional distinction between sex and gender. Gender is defined relationally: Men and women are two groups defined by their social relations to each other. I've argued elsewhere,[22] drawing on the work of Catharine MacKinnon, that we can usefully model one process by which gender is constructed roughly as follows: The ideal of Woman is an externalization of men's desire (so-called Woman's Nature is what men find desirable); this ideal is projected onto individual females and is regarded as intrinsic and essential to them. Accepting these attributions of Womanhood, individual women then internalize the norms appropriate to the ideal and aim to conform their behavior to them; and, in general, behavior towards women is "justified" by reference to this ideal. This, in turn, is responsible for significant empirical differences between men and women.

In this example, individual women are *discursively constructed*; that is, we are the individuals we are because of the attribution (and self-attribution) of Womanhood to us or, more simply, because we've been viewed (and so treated) as having a Woman's Nature. Because discursive construction is a kind of causal construction, it is also correct to say that individual women are *causally constructed*. The ideal of Woman's Nature, however, is *strongly pragmatically constructed*; it is an illusion projected onto women whose basis lies in complex social-sexual relations, not in the intrinsic or essential features of women. As in the case of "cool," we debunk the idea of Woman's Nature and find two concepts at work: The manifest concept of Woman's Nature—understood as defining *what women are by nature* in traditional terms—is an illusion; the operative concept being masked by it is constitutively constructed in terms of men's (socially conditioned) sexual

[22] Haslanger, "On Being Objective and Being Objectified." Chapter 1 of this volume.

responses. Further, the distinctions between both man and woman, and male and female (taken as groups of individuals) are *weakly pragmatically constructed*; the fact that we draw these distinctions as we do is to be at least partly explained by social factors, though there are also very real differences between both men and women, male and female.

To summarize, the following would be plausible examples of each kind of construction:

> *Discursively (and so causally) constructed*: individual women; cool dudes.
> *Strongly pragmatically constructed*: Woman's Nature; intrinsic coolness.
> *Constitutively constructed*: the operative concept of "coolness"; the operative concept of "Woman's Nature."
> *Weakly pragmatically constructed*: The distinction between men and women, between male and female; the distinction between those who wear black t-shirts more than once a week and those who don't.

3. The Social Construction of Reality

Given the different kinds of social construction just sketched, there are a variety of different senses we might give to the claim that *reality* is socially constructed. For example, the claim might be that human beings are in some significant way involved in bringing about or constituting everything there is or, more specifically, that our linguistic and conceptual activities are responsible for how things are. Alternatively, the claim might be that how we conceive of reality is determined wholly or partly by social factors. Are any of these claims plausible? And if so, should we be led to give up the idea that there is a world (in some sense) "independent" of us?

The Causal Construction of Reality

Consider, for the moment, causal construction. Is it plausible that the entire world—not just Earth, but everything there is—is a human artifact, even allowing that the mechanisms of construction might be highly complex and mediated? Is the world, for example, a product of our efforts at classification? I don't think so. Clearly human beings have had an enormous effect on things: Mountains are damaged by acid rain, the polar icecaps are melting. And it is equally clear that our actions and classificatory efforts can make a big difference to the nonhuman world: Microbes adapt to our classifications of them by becoming immune to our antibacterial agents.[23] But not everything is so responsive to our

[23] See Ian Hacking, "Making Up People," in Thomas Heller, Morton Sosna, and David Wellberg, eds., *Reconstructing Individualism* (Stanford, Calif.: Stanford University Press, 1986), 222–36.